

Intro to Deep Learning with Keras

Thu 20th September 2018 5pm - 6:30pm



Install

pip install tensorflow

pip install keras

Then visit

coefficient.training

ARTIFICIAL NEURAL NETWORKS

1943: The Mechanical Tortoise

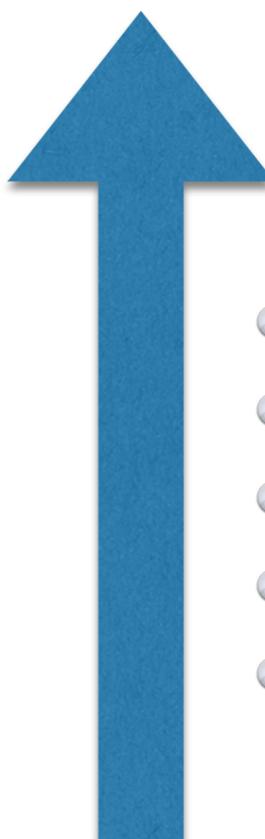


1956: The Dartmouth Conference

“ Every aspect of learning or other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it ”

Top Down?

- Make rules
- Wait for something to break
- Refine the rules
- "Expert systems"
- Brittle & expensive



- Define goals
- Systems have complex design
- Systems find patterns
- Lots of computing power
- More like nature

Bottom Up?

1970s: The AI Winter

“ In **three to eight years** we will have a machine with
the **general intelligence** of an **average human being** ”

– 1970, Marvin Minsky, *Founder of MIT Artificial Intelligence Lab*

“ There has been **no machine translation** of general scientific
text and **none is in immediate prospect** ”

– 1966 report by US government advisory committee

“ **In no part of the field** have discoveries made
so far produced the **major impact that was promised** ”

– 1973, Professor Sir James Lighthill, "Artificial Intelligence: A General Survey"

80s - 90s: Smaller Ambitions, Big Business



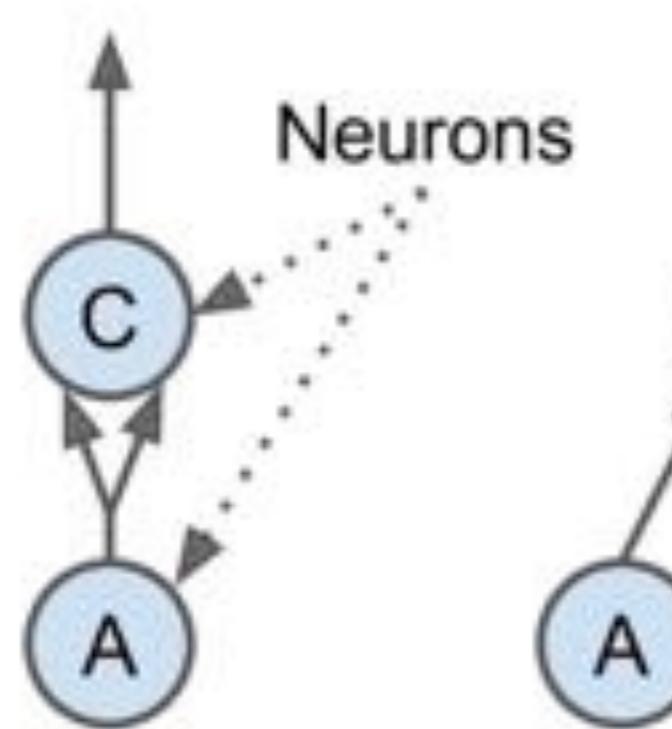
1981:
Expert systems at **DEC** configure
orders for new computer systems,
saves company \$40m per year



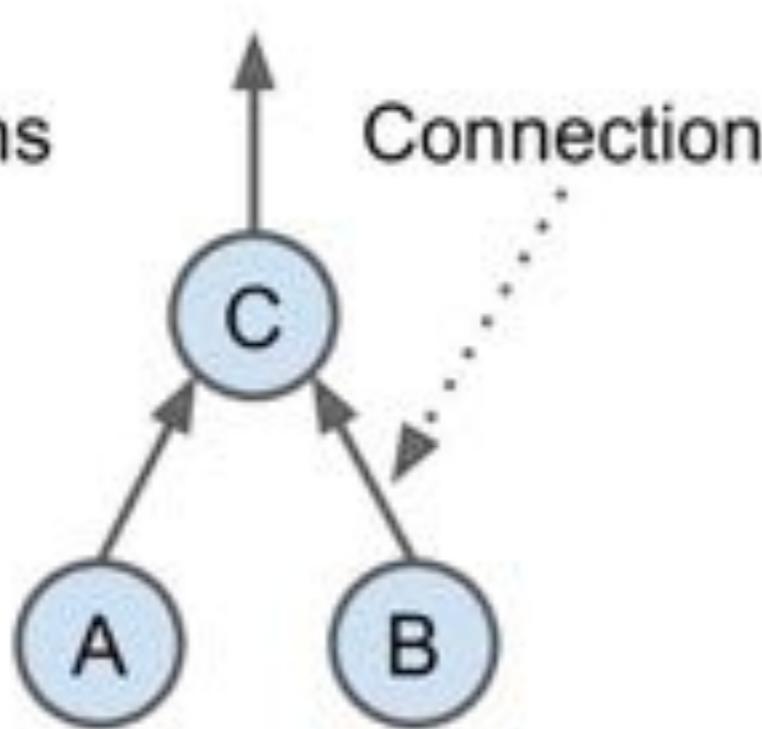
1997:
IBM's Deep Blue defeats
world chess champion
Gary Kasparov

Artificial Neurons

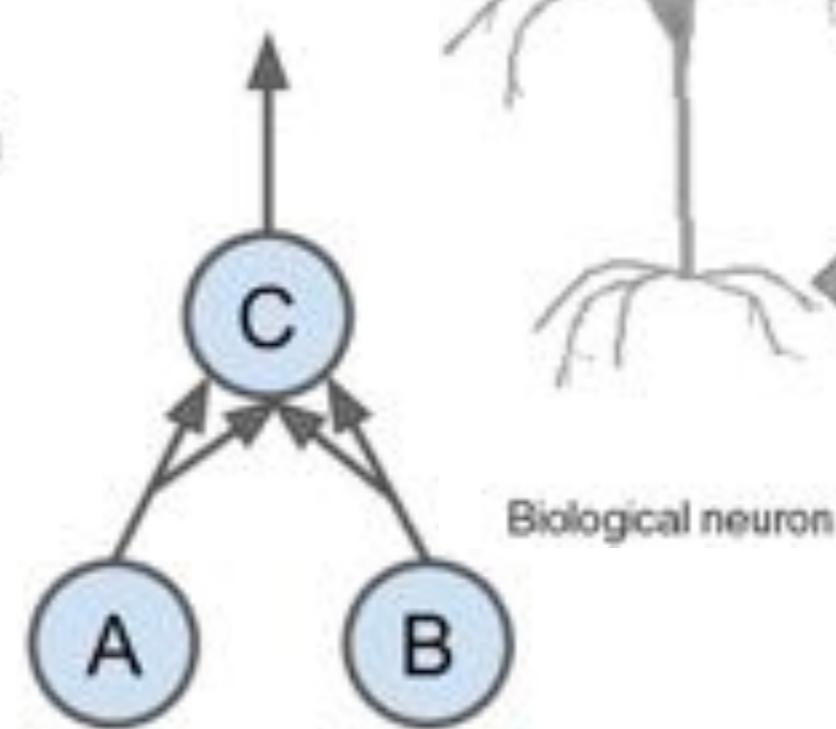
- ▶ Proposed by *McCulloch & Pitts*
- ▶ Inspired by biological neurons, artificial neurons can replicate logical operations using **binary I/O** and **thresholds**.
- ▶ e.g. neuron fires if **2+ inputs fire**.



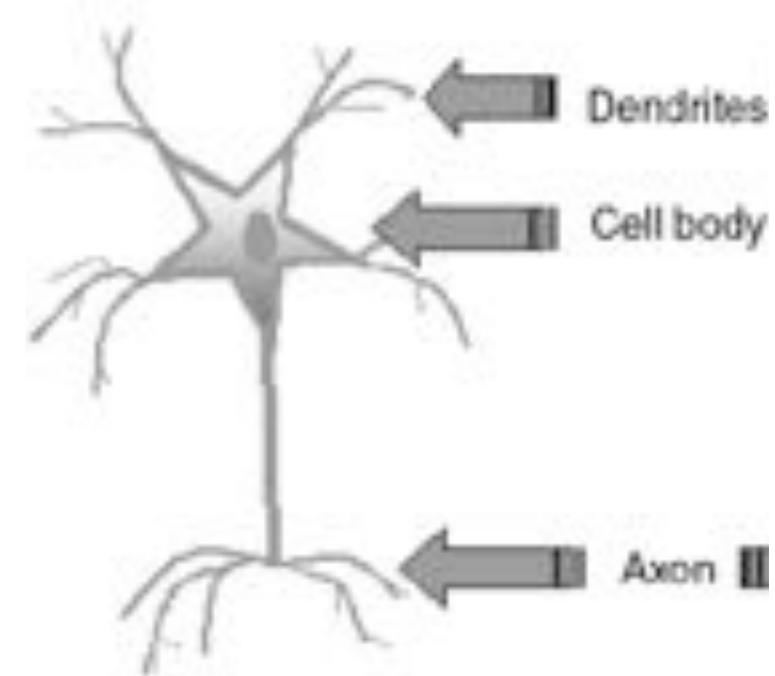
$$C = A$$



$$C = A \wedge B$$



$$C = A \vee B$$

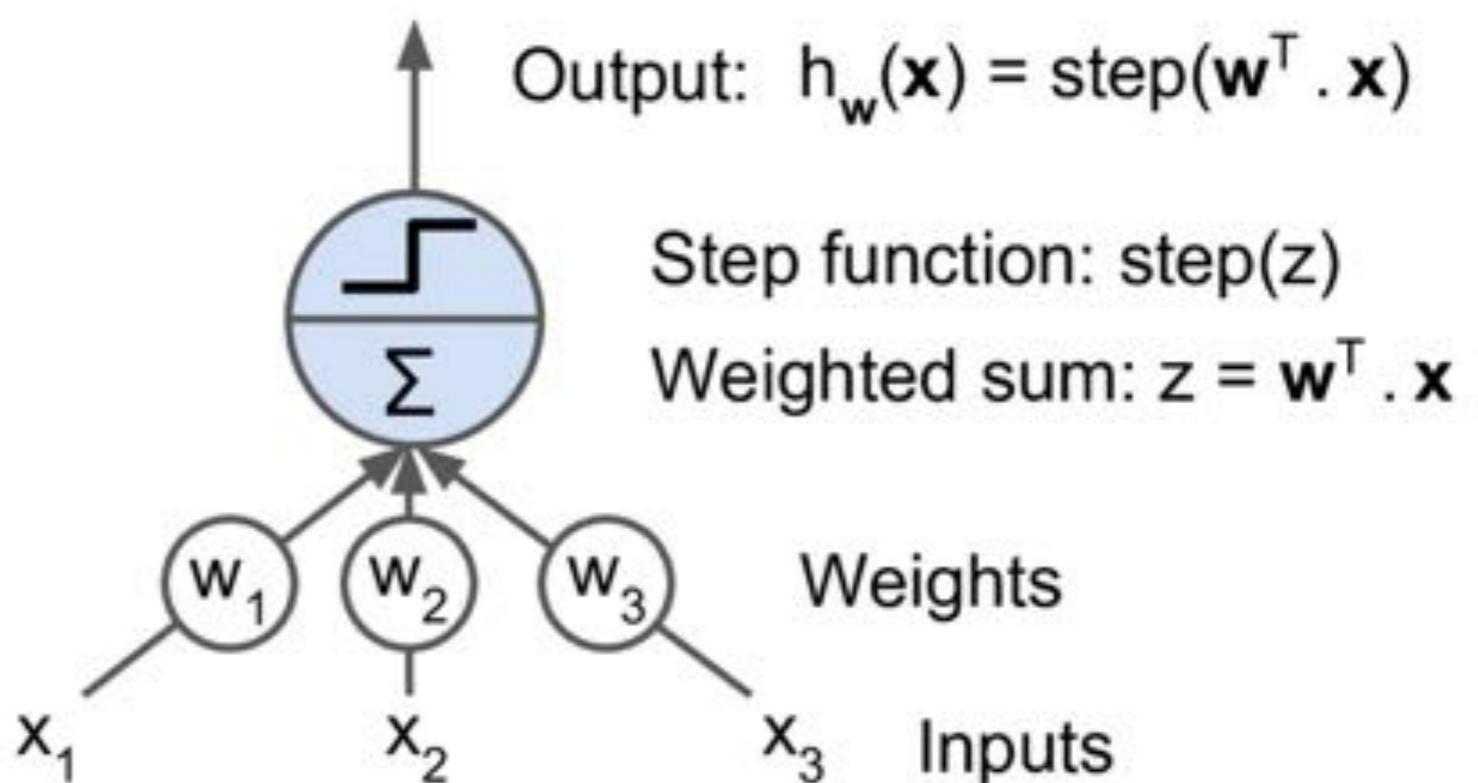


The Linear Threshold Unit

- ▶ Frank Rosenblatt, 1957
- ▶ Replaces neuron with **linear threshold unit (LTU)**
 - ▶ Binary I/O → **continuous I/O**
 - ▶ 1. Compute **weighted sum of inputs** (linear, i.e. dot product)
 - ▶ 2. Apply **step function** (e.g. *heaviside* or *sign*)
- ▶ Can be used as a simple binary classifier (add a bias + sigmoid step function and you have logistic regression!)

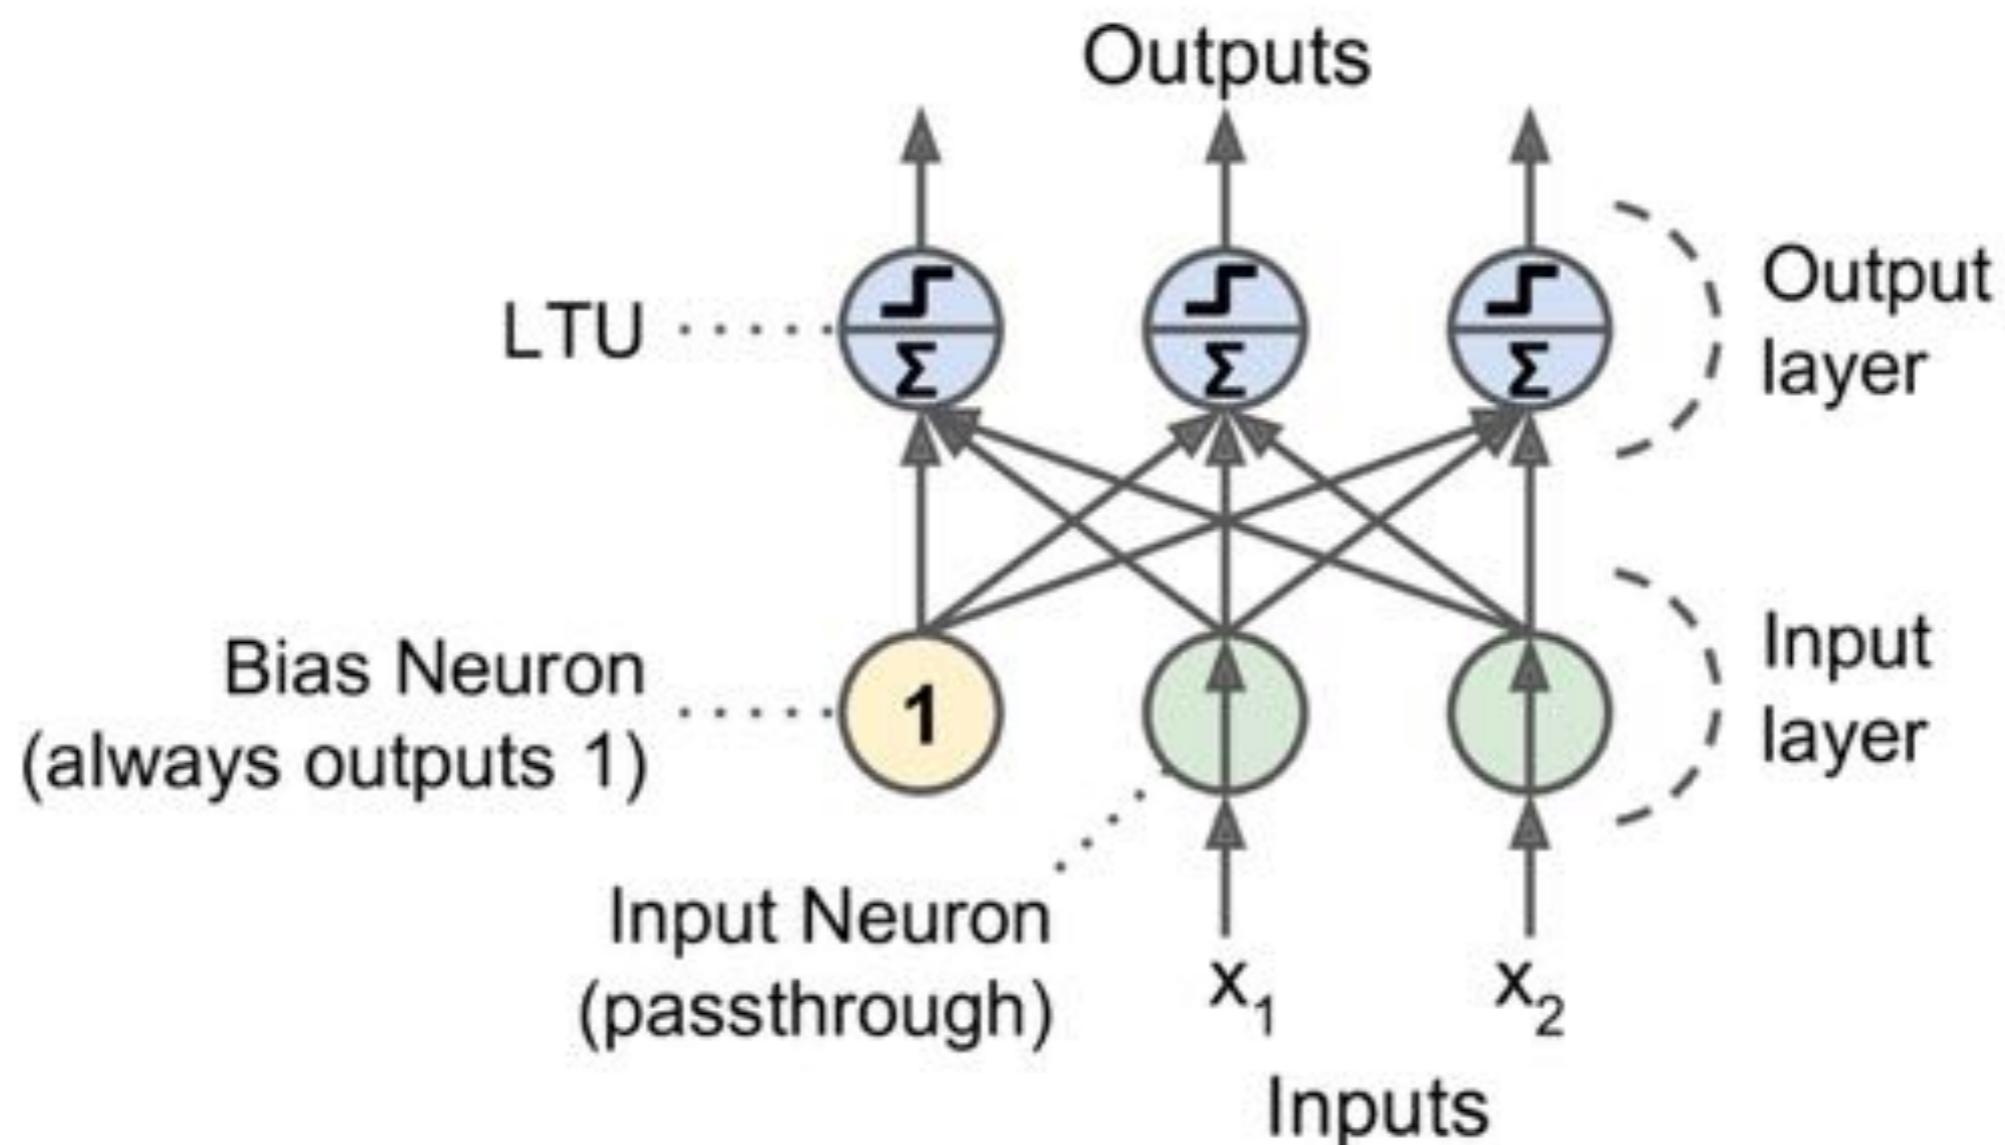
$$\text{heaviside}(z) = \begin{cases} 0 & \text{if } z < 0 \\ 1 & \text{if } z \geq 0 \end{cases}$$

$$\text{sgn}(z) = \begin{cases} -1 & \text{if } z < 0 \\ 0 & \text{if } z = 0 \\ +1 & \text{if } z > 0 \end{cases}$$



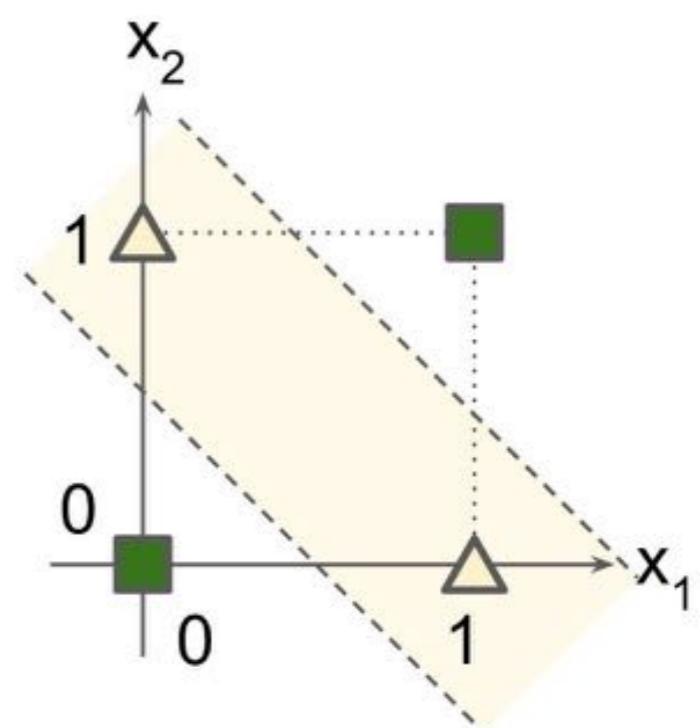
The Perceptron

- ▶ A Perceptron is simply a **layer** of LTUs plus a **bias** term.
- ▶ Every input feature flows into every LTU.
- ▶ The Perceptron shown below can model three-class binary classification problems.



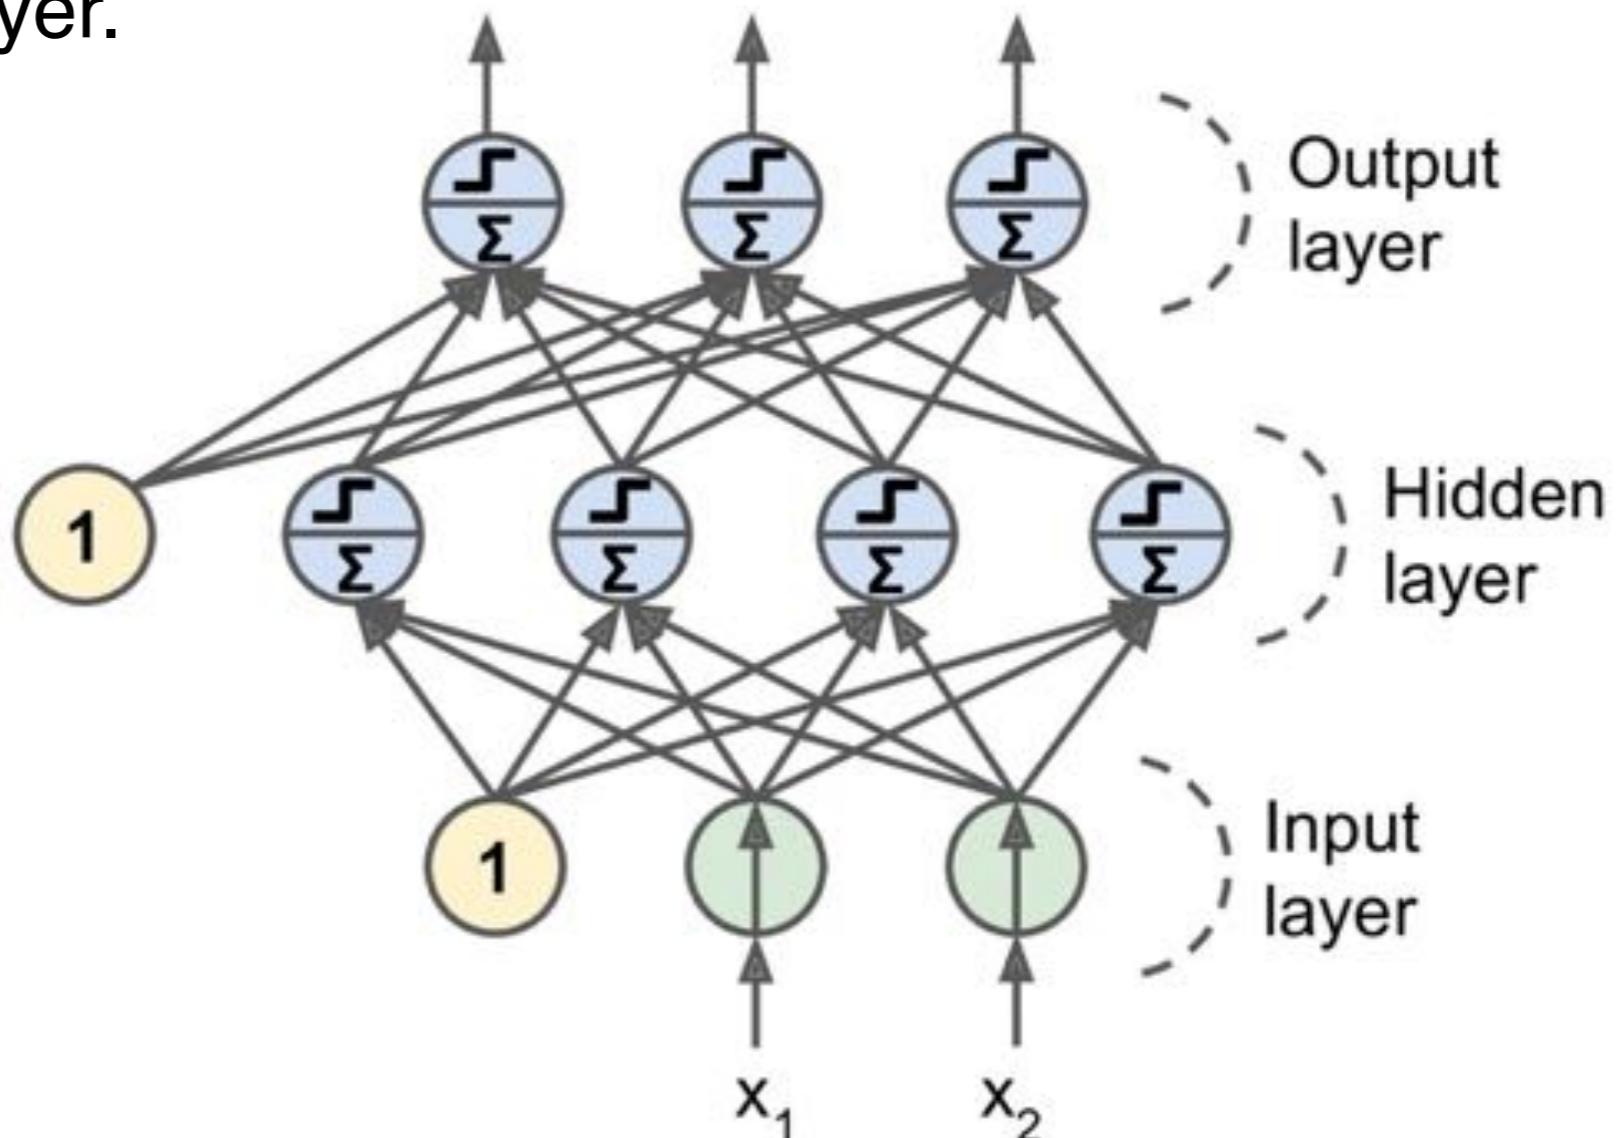
The Perceptron

- ▶ **Training Perceptrons**
 - ▶ "Cells that fire together, wire together" – Siegrid Löwel
 - ▶ Input single training instances one at a time. For each incorrect answer, reinforce the weights that would've helped get it right.
- ▶ Same as the SGDClassifier in sklearn with **loss='perception'**, **learning_rate='constant'**, **eta0=1**, **penalty=None**.
- ▶ Similar to Multi-class Logistic Regression, except that logistic regression is able to generate probabilities (via sigmoid function).
- ▶ **Issue:** Cannot solve XOR ("exclusive OR") —> (*Minsky & Papert, 1969; cue disappoint & AI Winter*)



The Multi-Layer Perceptron (MLP)

- ▶ A Multi-Layer Perceptron is a stack of multiple Perceptrons:
 - ▶ One input ("passthrough") layer.
 - ▶ At least one layer of "hidden" LTUs. If there are two or more hidden layers this is called a "*deep neural network*" (DNN).
 - ▶ One final output layer.



Activation Functions

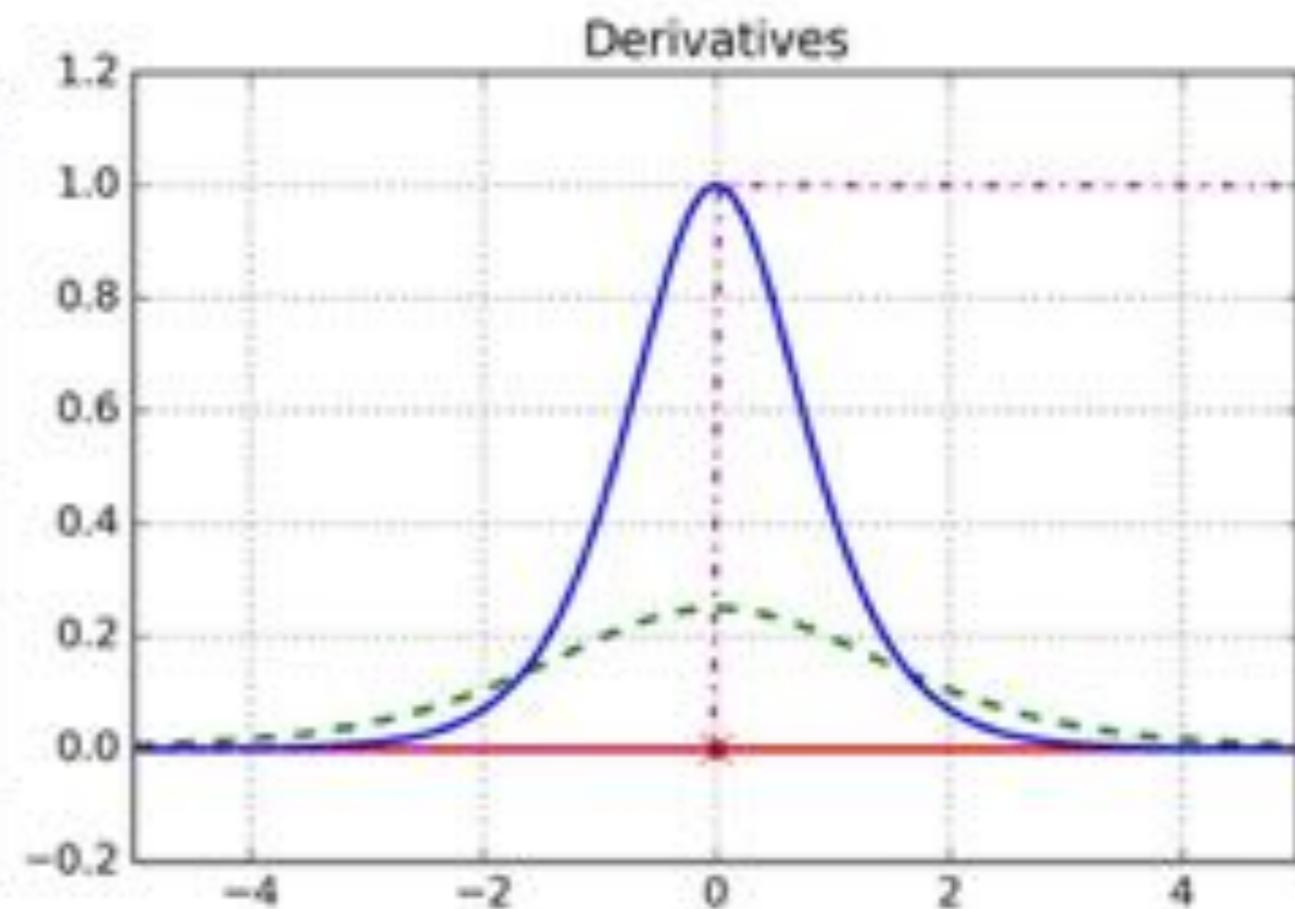
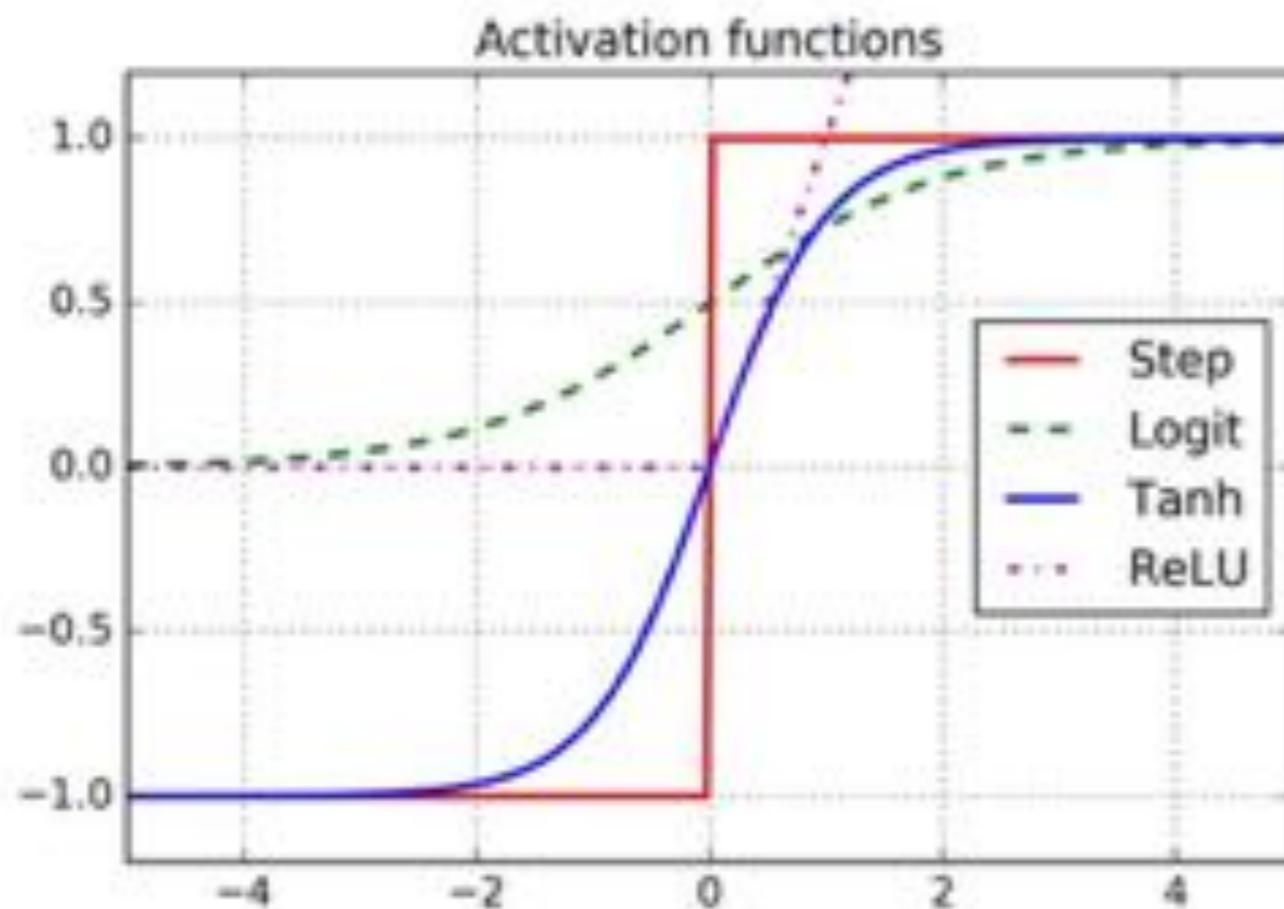
In order to calculate gradients, we need a smooth differentiable function (i.e. not a step function!).

Sigmoid was common early on as nature uses this in biological neurons.

$$\text{sigmoid: } \sigma(x) = \frac{1}{1+e^{-x}}$$

$$\tanh: \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$RELU: \text{relu}(x) = \max(0, x)$$





LAB:

MULTI-LAYER

PERCEPTRON



Keras



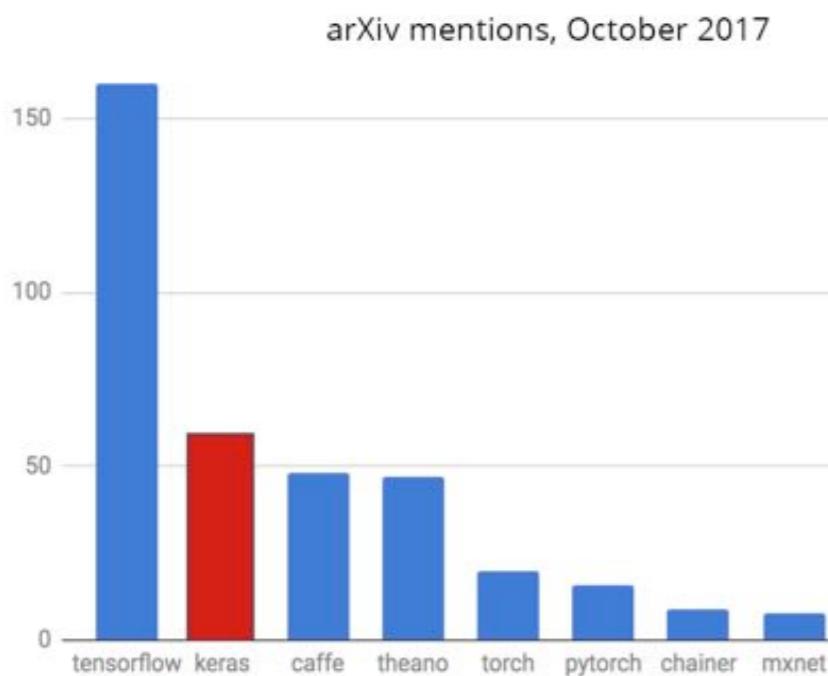
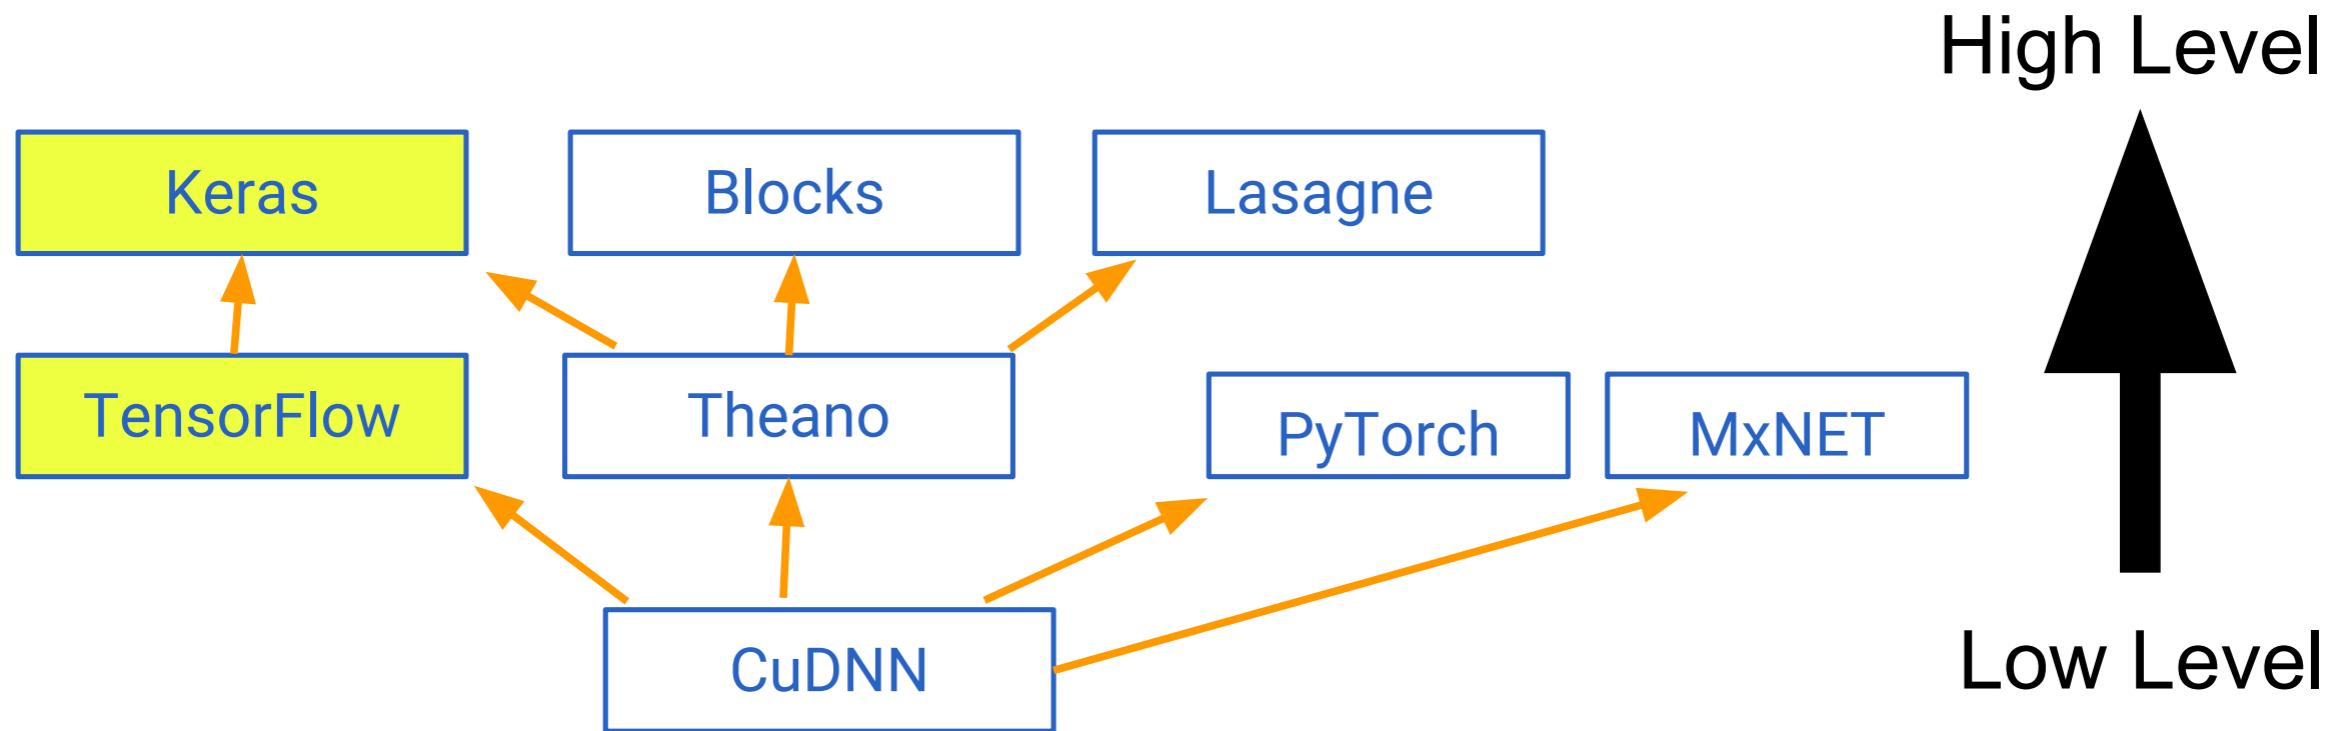
"Keras was developed with a focus on enabling fast experimentation. Being able to go from idea to result with the least possible delay is key to doing good research."

► Features

- High-level neural network API written in Python.
- Just an API! Doesn't lock you into a framework, can currently run on top of TensorFlow (deeply integrated), CNTK (Microsoft) or Theano, plus MXNet (Apache/Amazon) coming...
- Supports Feed-Forward Neural Networks, Convolutional NNs, Recurrent NNs, and complex combinations of all the above.
- Runs seamlessly on both CPU and GPU.

► Workflow: **Define → Compile → Fit → Predict**

The Deep Learning Landscape



Why has Keras been so successful lately at Kaggle competitions?

< 200k users (Nov'17) & #2 for scientific mentions

Keras runs everywhere!

► Run Keras on...

- ▶ **iOS** via CoreML (official Keras support provided by Apple)
- ▶ **Android** via TensorFlow Android runtime.
- ▶ **Browsers** via GPU-accelerated JS runtimes (keras.js, WebDNN)
- ▶ **Google Cloud** via TensorFlow-Serving
- ▶ **JVM** via DL4J model import functionality

► Supported hardware platforms

- ▶ **NVIDIA GPUs** via cuDNN
- ▶ **Google TPUs** ("Tensor Processing Units", each providing up to 180 teraflops of performance)
- ▶ **OpenCL-enabled GPUs** (including macOS!) via PlaidML's Keras backend
- ▶ Built-in **multi-GPU support** plus **distributed training** via **Horovod** (Uber), **Dist-Keras** (Spark/CERN), and **Elephas**.



LAB:

BUILD A MLP IN KERAS



LAB:

CONVOLUTIONAL

NEURAL NETWORKS

IN KERAS